

Harsh Lara

☎ 412-583-1116 | 🌐 harshlara | 🏠 harshlara.github.io | ✉ hlara.tech@gmail.com

Work Experience

Resolve AI

San Francisco, CA

Sept 2025 – Present

MEMBER OF TECHNICAL STAFF

- Among Resolve AI's early technical hires, scaling the company from Seed (\$115M valuation) to [Series A \(\\$1.5B\)](#) with backing from Greylock, Lightspeed, DST Global, and Jeff Dean
- Architected and shipped Resolve AI's mitigations system: autonomous code mitigation and remediation agents that resolve production incidents end-to-end, reducing customer incident mean-time-to-resolve (MTTR)
- Drove launch of product features that onboarded enterprise customers including Coinbase, DoorDash, and MongoDB
- Led evaluation strategy across Resolve AI's agentic incident-resolution stack including investigation, root cause analysis, mitigation accuracy, and end-to-end agent reliability
- Represented Resolve AI at [HumanX 2026](#)

Google DeepMind

Mountain View, CA

Sept 2021 – Sept 2025

RESEARCH ENGINEER, GEMINI

- Led post-training for improving tool-use capabilities in [Project Astra](#) and [Gemini Live](#), enabling critical integration with Google Workspace and significantly enhancing productivity for users
- Co-led the design and implementation of [MoDE](#): A multi-task parameter efficient fine-tuning technique, outperforming existing methods and achieving significant performance improvements across over 700 tasks ([published at NAACL 2025](#))
- Core contributor to [OmegaPRM](#): A novel MCTS algorithm for automated process supervision, directly improving mathematical reasoning in Gemini models from 51% to 69.4% on MATH500 and 86.4% to 93.6% on GSM8K
- Core contributor to [RecLLM](#): LLM-based Conversational Recommender System (CRS) for YouTube videos using LaMDA; designed a controllable LLM-based user simulator to address data scarcity in CRS training and evaluation

Microsoft

Bellevue, WA

Dec 2019 – Sept 2021

APPLIED SCIENTIST, AZURE MACHINE LEARNING

- Drove AutoML product strategy and delivery for Azure Machine Learning, shipping Classification, Regression, and Data Labeling capabilities as managed cloud services (ML-as-a-Service)
- Served as Tech Lead and Applied Sciences Lead across multiple AutoML projects, owning end-to-end design and delivery from research to production
- Led design and implementation of end-to-end neural network infrastructure that significantly improved training and inference efficiency for text data-labeling pipelines

Apple

Cupertino, CA

Summer 2019

MACHINE LEARNING ENGINEER INTERN, SIRI

- Developed an innovative method to dynamically interpolate probabilistic language-models on-device achieving multi-domain increase (>15%) in ASR accuracy. Received strong recognition from senior AI/ML leads.

Education

Carnegie Mellon University, School of Computer Science

Pittsburgh, PA

MS, COMPUTATIONAL DATA SCIENCE

- Selected courses: Cloud Computing | Computer Systems | Intro to Machine Learning | Language and Statistics | Large-Scale Multimedia Analysis | Applied Machine Learning | Deep Learning | Computational Ethics for NLP

National Institute of Technology Karnataka (NITK)

Surathkal, India

B TECH, COMPUTER SCIENCE AND ENGINEERING

- Selected courses: Operating Systems | Databases | Distributed Systems | Object Oriented Programming | Advanced Algorithms | Data Structures | Image Processing | Data Communication | Graph Theory | Discrete Math | Linear Algebra

Skills

Focus Areas: Generative AI, Agentic AI Systems, Evals, Harnesses, Post-training LLMs, LLMs, Artificial Intelligence
Languages/Libraries: Python, TypeScript, Go, C/C++, Java, SQL
Tools/Technologies: PyTorch, TensorFlow, Clickhouse, Grafana, GCP, Azure, AWS

Publications

- [MoDE: Effective Multi-task Parameter Efficient Fine-Tuning with a Mixture of Dyadic Experts](#), NAACL 2025
- [Improve Mathematical Reasoning in Language Models by Automated Process Supervision](#), arXiv 2024
- [Leveraging Large Language Models in Conversational Recommender Systems](#), arXiv 2023
- [Evaluation of Synthetic Datasets for Conversational Recommender Systems](#), NeurIPS 2022

Patents

- [Generating Synthetic Conversation Data for Training and Evaluation of Natural Language Models](#) (Google)
- [Neural Networks with Effective Parameter Efficient Adaptations](#) (Google)
- [Efficient Decoding of Output Sequences Using Early Exit Adaptors](#) (Google)

Honors & Awards

- Indian Academy of Sciences (IAS) Summer Research Fellowship 2016 — selected as 1 of ~1.5K Fellows from 100K+ applicants; researched Named Entity Recognition and Relation Extraction for the BioNLP 2016 Competition at IIT Guwahati